Oracle for Developers
(DBMS SQL)
Lab Book

Document Revision History

Date	Revision No.	Author	Summary of Changes
05-Feb-2009	0.1D	Rajita Dhumal	Content Creation
09-Feb-2009		CLS team	Review
20-May-2011	2.0	Sathiabama R	Integration Refinements
23-May-2011	2.0	Anu Mitra	Review

Table of Contents

Getting Started

Overview

This lab book is a guided tour for learning DBMS SQL. It comprises 'To Do' assignments. Follow the steps provided and work out the 'To Do' assignments.

Setup Checklist for DBMS SQL

Here is what is expected on your machine in order for the lab to work.

Minimum System Requirements

- Intel Pentium 90 or higher (P166 recommended)
- Microsoft Windows 95, 98, or NT 4.0, 2k, XP.
- Memory: 32MB of RAM (64MB or more recommended)

Please ensure that the following is done:

- Oracle Client installed.
- Connectivity to Oracle Server

Instructions

- For all coding standards refer Appendix A. All lab assignments should refer coding standards.
- Create a directory by your name in drive <drive>. In this directory, create a subdirectory <DBMS SQL>_assgn. For each lab exercise create a directory as lab <lab number>.

Learning More (Bibliography if applicable)

NA

©2016 Capgemini. All rights reserved.

The information contained in this document is proprietary and confidential. For Capgemini only. | **4** / 29

Problem Statement / Case Study

Table descriptions

Designation_Masters

Name	Null?	Type
Design_code	Not Null	Number(3)
Design_name		Varchar2(50)

Department_Masters

Name	Null?	Type
Dept_Code	Not Null	Number(2)
Dept_name		Varchar2(50)

Student_Masters Table

Name	Null?	Type
Student_Code	Not Null	Number(6)
Student_name	Not Null	Varchar2(50)
Dept_Code		Number(2)
Student_dob		Date
Student_Address		Varchar2(240)

Student_Marks

Name	Null?	Type
Student_Code		Number(6)
Student_Year	Not Null	Number
Subject1		Number(3)
Subject2		Number(3)
Subject3		Number(3)

Staff_Masters

Name	Null?	Type
Staff_code	Not Null	Number(8)
Staff_Name	Not Null	Varchar2(50)
Design_code		Number

Dept_code		Number
HireDate		Date
Staff_dob		Date
Staff_address		Varchar2(240)
Mgr_code		Number(8)
Staff_sal		Number (10,2)

Book_Masters

Name	Null?	Type
Book_Code	Not Null	Number(10)
Book_Name	Not Null	Varchar2(50)
Book_pub_year		Number
Book_pub_author	Not Null	Varchar2(50)

Book_Transactions

Name	Null?	Type
Book_Code		Number
Student_code		Number
Staff_code		Number
Book_Issue_date	Not Null	Date
Book_expected_return_date	Not Null	Date
Book_actual_return_date		Date

|

Lab 1. Data Query Language

Goals	• Query the database . • Usage of various operators in select statements. • NVL Function.
Time	1 hr 30 min

1.1: Data Query Language

1. Retrieve the details (Name, Salary and dept code) of the employees who are working in department 20, 30 and 40.

2. List the details of the employees with user defined Column headers.

3. Display the code, subjects and total marks for every student. Total Marks will be calculated as Subject1+Subject2+Subject3. (Refer Student_Marks table)

4. List the details of the staff whose designations are either PROFESSOR or LECTURER.

5. List the code, name, and department number of the employees who have experience of more than 18 years.

6. List the name and Designations of the staff who have joined before Jan 2003.

7. List the name, designation, and income for 10 years of the employees who are working in departments 10 and 30.

8. List the name and experience (in years) of employees who are working as LECTURER.

9. Display name concatenated with dept code separated by comma and space. Name the column as 'Student Info'.

10. List the Name and Salary of the staff who are earning between 12000 and 25000. Sort them based on their salaries and name.

11. Display employees who do not have manager.

12. Write a query which will display name, department code and date of birth of all students who were born between January 1, 1981 and March 31, 1983. Sort it based on date of birth (ascending).

13. Get the Department number, and sum of Salary of all non managers where the sum is greater than 20000.

14. Display the details of books that have not been returned and expected return date was last Monday.

15. Display the name and department code of students. If student does not belong to any department, display "No Department". Label the column as "Department". (Hint: Use NVL function)

16. Display the name and salary of the staff. Salary should be represented as X. Each X represents a 1000 in salary.

Sample Output

JOHN	10000	XXXXXXXXXX
ALLEN	12000	XXXXXXXXXXXX

|

Lab 2.   Single Row Functions

Goals	• Querying tables using single row functions • Querying tables using date functions • Querying tables using number functions • Querying tables using group functions
Time	1 hr 45 min

2.1: Single Row Functions:

1. Display name and date of birth of students where date of birth must be displayed in the format similar to "January, 12 1981" for those who were born on Saturday or Sunday.

2. Display each staff name and number of months they worked for the organization. Label the column as 'Months Worked'. Order your result by number of months employed. Round the number of months to closest whole number.

3. List the details of the employees, whose names start with 'A' and end with 'S'.

4. List the name and job of the employees whose names should contain N as the second or third character, and ending with either 'N' or 'S'.

5. Create a query which will display Staff Name, Salary of each staff. Format the salary to be 15 character long and left padded with '$'.

6. List the names of the Employees having '_' character in their name.

7. List the details of the employees who have joined in December (irrespective of the year).

8. Write a query that displays Staff Name, Salary, and Grade of all staff. Grade depends on the following table.

Salary	Grade
Salary >=50000	A

Salary >= 25000 < 50000	B
Salary>=10000 < 25000	C
OTHERS	D

9. Display the Staff Name, Hire date and day of the week on which staff was hired. Label the column as DAY. Order the result by the day of the week starting with Monday.

10. Show staff names with the respective numbers of asterisk from Staff table.

11. Show staff names with the respective numbers of asterisk from Staff table except first and last characters. For example: KING will be replaced with K**G.

12. Show all staffs who were hired in the first half of the month.

13. Display the staff name, hire date and day of the week on which the staff joined. Order the results by the day of the week starting with Monday.

14. Write a query to find the position of third occurrence of 'i' in the given word 'Mississippi'.

15. Write a query to find the pay date for the month. Pay date is the last Friday of the month. Display the date in the format "Twenty Eighth of January, 2002". Label the heading as PAY DATE.

2.2: Group Functions:

16. Display the Highest, Lowest, Total & Average salary of all staff. Label the columns Maximum, Minimum, Total and Average respectively. Round the result to nearest whole number.

17. Edit the above query and display the same for each Department Name.

18. Write a query to display number of people in each Department. Output should display Department Code, Department Name and Number of People.

©2016 Capgemini. All rights reserved.

The information contained in this document is proprietary and confidential. For Capgemini only. | **10** / 29

19. Determine the number of managers without listing them. Label the column as 'Total Number of Managers'.

20. Display Manager Code, Manager Name and salary of lowest paid staff in that manager's team. Exclude any group where minimum salary is less than 10000. Order the result on descending order of salary.

©2016 Capgemini. All rights reserved.

The information contained in this document is proprietary and confidential. For Capgemini only. | **11** / 29

Lab 3.   JOINS AND SUBQUERIES

Goals	• Querying multiple tables using joins • Querying tables using subqueries • Querying tables using set operators
Time	2 hr 30 min

3.1: Joins and Subqueries

1. Write a query which displays Staff Name, Department Code, Department Name, and Salary for all staff who earns more than 20000.

2. Write a query to display Staff Name, Department Code, and Department Name for all staff who do not work in Department code 10 and have 'A' in their name.

3. Display Staff Code, Staff Name, Department Name, and his manager's number and name. Label the columns Staff#, Staff, Mgr#, Manager.

4. Create a query that will display Student Code, Student Name, Department Name, Subjec1, Subject2, and Subject3 for all students who are getting 60 and above in each subject from department 10 and 20.

5. Create a query that will display Student Code, Student Name, Book Code, and Book Name for all students whose expected book return date is today.

6. Create a query that will display Staff Code, Staff Name, Department Name, Designation, Book Code, Book Name, and Issue Date. For only those staff who have taken any book in last 30 days.

7. Generate a report which contains the following information.
Staff Code   Staff Name   Designation   Department Name
Department Head
For all staff excluding HOD (List should not contain the details of Department head).

8. Generate a report which contains the following information

Student Code   Student Name Department Name Total Marks
 HOD Name
Sort the output on Department Name and Total Marks.

9. Generate a report which contains the following information.

Staff Code, Staff Name, Designation, Department, Book Code, Book Name,
 Author, Fine
For the staff who have not return the book. Fine will be calculated as Rs. 5 per day.
Fine = 5 * (No. of days = Current Date – Expected return date).

10. List Staff Code, Staff Name, and Salary for those who are getting less than the average salary of organization.

11. List the Staff Code, Staff Name who are not Manager.

12. Display Author Name, Book Name for those authors who wrote more than one book.

13. Display Staff Code, Staff Name, and Department Name for those who have taken more than one book.

14. Display top ten students for a specified department. Details are:
 Student Code, Student Name, Department Name, Subject1, Subject2,
 Subject3, Total.

15. Display the Staff Name, Department Name, and Salary for those staff who are getting less than average salary in their own department

16. Create a query that will display the Staff Name, Department Name, and all the staff that work in the same department as a given staff. Give the column as appropriate label.

17. List the Student Code, Student Name for that student who got highest marks in all three subjects in Computer Science department for current year.

©2016 Capgemini. All rights reserved.

The information contained in this document is proprietary and confidential. For Capgemini only. | **13** / 29

18. Display the Student Code, Student Name, and Department Name for that department in which there are maximum number of student are studying.

19. Display Staff Code, Staff Name, Department Name, and Designation for those who have joined most recently.

20. Display the Manager Name and the total strength of his/her team.

3.2: Set Operators

Observe the following queries and do the given assignments.

```
create table previous_products (
      pid int,
      name varchar(40),
      unit varchar(40),
      price int,
      stock int );
```

Example 1: Sample Code

```
create table current_products (
      pid int,
      name varchar(40),
      unit varchar(40),
      price int,
      stock int );
```

Example 2: Sample Code

```
insert into previous_products values(7,'Uncle Bob''s Organic Dried Pears','12 - 1 lb
pkgs.',30.00,15);
insert into previous_products values(8,'Northwoods Cranberry Sauce','12 - 12 oz
jars',40.00,6);
insert into previous_products values(1,'Chang','24 - 12 oz bottles',19.00,17);
insert into previous_products values(3,'Aniseed Syrup','12 - 550 ml
bottles',10.00,13);
insert into previous_products values(4,'Chef Anton''s Cajun Seasoning','48 - 6 oz
jars',22.00,53);
insert into previous_products values(5,'Chef Anton''s Gumbo Mix','36
boxes',21.35,0);
```

©2016 Capgemini. All rights reserved.

The information contained in this document is proprietary and confidential. For Capgemini only. | **14** / 29

```
insert into previous_products values(6,'Grandma''s Boysenberry Spread','12 - 8 oz
jars',25.00,120);
```

Example 3: Sample Code

```
insert into current_products values(7,'Uncle Bob''s Organic Dried Pears','12 - 1 lb
pkgs.',30.00,15);
insert into current_products values(8,'Northwoods Cranberry Sauce','12 - 12 oz
jars',40.00,6);
insert into current_products values(9,'Mishi Kobe Niku','18 - 500 g pkgs.',97.00,29);
insert into current_products values(10,'Ikura','12 - 200 ml jars',31.00,31);
insert into current_products values(11,'Queso Cabrales','1 kg pkg.',21.00,22);
insert into current_products values(5,'Chef Anton''s Gumbo Mix','36 boxes',21.35,0);
insert into current_products values(6,'Grandma''s Boysenberry Spread','12 - 8 oz
jars',25.00,120);
```

Example 4: Sample Code

1. Get the details of all products irrespective of the fact whether they are in previous set or current set.

2. Get the details of all products along with the repetition of those that were present both in the previous and current sets.

3. Get the details of only those products which were present in the previous set and are still continuing.

4. Get the details of all obsolete products (no longer continued).

Lab 4.   Database Objects

Goals	Following set of questions are designed to implement the following concepts • Creating Database objects like tables, views , etc • Modifying Database objects • Deleting Database objects • Usage of Data Dictionary tables
Time	4 hr 30 min

4.1: Database Objects

1. Create the Customer table with the following columns.

      Customerid   Number(5)

      CustomerName Number(10)

      Address1   Varchar2(30)

      Address2   Varchar2(30)

2. Modify the Customer table CustomerName column of datatype with Varchar2(30). CustomerName should not accept Nulls.

3. Add the following Columns to the Customer table.

      Gender   Varchar2(1)

      Age   Number(3)

      PhoneNo   Number(10)

4. Insert rows with the following data in to the Customer table.

Insert into customer values: (1000, 'Allen', '#115 Chicago', '#115 Chicago', 'M', '25, 7878776')

In similar manner, add the below records to the Customer table:

* 1000, Allen, #115 Chicago, #115 Chicago, M, 25, 7878776

* 1001, George, #116 France, #116 France, M, 25, 434524

* 1002, Becker, #114 New York, #114 New York, M, 45, 431525

5. Add the Primary key constraint for CustomerId with the name CustId_Prim.

6. Insert the row given below in the Customer table and see the message generated by the Oracle server.

1002, John, #114 Chicago, #114 Chicago, M, 45, 439525

7. Disable the constraint on CustomerId, and insert the following data:

- 1002, Becker, #114 New York, #114 New york , M, 45, 431525
- 1003, Nanapatekar, #115 India, #115 India , M, 45, 431525

8. Enable the constraint on CustomerId of the Customer table, and see the message generated by the Oracle server.

9. Drop the constraint CustId_Prim on CustomerId and insert the following Data. Alter Customer table, drop constraint Custid_Prim.

- 1002, Becker, #114 New York, #114 New york , M, 45, 431525, 15000.50
- 1003, Nanapatekar, #115 India, #115 India , M, 45, 431525, 20000.50

10. Delete all the existing rows from Customer table, and let the structure remain itself using TRUNCATE statement.

11. In the Customer table, add a column E_mail.

12. Drop the E_mail column from Customer table.

13. Add a new column EmailId to Customer table.

14. Mark EmailId column as unused before dropping it.

15. Drop the unused EmailId column from the Customer table.

16. Define the COMMENT 'Customers Details' for Customer table.

17. Use Data Dictionary USER_TAB_COMMENTS to view the created comment.

18. Define the COMMENT 'Personal Contact no' for the phoneno column of the Customer table.

19. Use Data Dictionary USER_COL_COMMENTS to view the created comment.

20. Create the Suppliers table based on the structure of the Customer table. Include only the CustomerId, CustomerName, Address1, Address2, and phoneno columns.

Name the columns in the new table as SuppID, SName, Addr1, Addr2, and Contactno respectively.

21. Drop the above table and recreate the following table with the name CustomerMaster.

CustomerId	Number(5) Primary key(Name of constraint is CustId_PK)
CustomerName	Varchar2(30) Not Null
Addressl	Varchar2(30) Not Null
Address2	Varchar2(30)
Gender	Varchar2(l)
Age	Number(3)
PhoneNo	Number(10)

22. Create the AccountsMaster table with the following Columns. Use sequence to generate Account number

CustomerId	Number(5)
AccountNumber	Number(10,2) Primary key(Name of constraint is Acc_PK)
AccountType	Char(3)
LedgerBalance	Number(10,2) Not Null

23. Relate AccountsMaster table and CustomerMaster table through CustomerId column with the constraint name Cust_acc.

24. Insert the following rows to the CustomerMaster table:

- 1000, Allen, #115 Chicago, #115 Chicago, M, 25, 7878776
- 1001, George, #116 France, #116 France, M, 25, 434524
- 1002, Becker, #114 New York, #114 New York, M, 45, 431525

25. Modify the AccountMaster table with the Check constraint to ensure AccountType should be either NRI or IND.

26. Insert 5 rows into the AccountsMaster table:

27. Modify the AccountsMaster table keeping a Check constraint with the name Balance_Check for the Minimum Balance which should be greater than 5000.

28. Modify the AccountsMaster table such that if Customer is deleted from Customer table then all his details should be deleted from AccountsMaster table.

29. Create Backup copy for the AccountsMaster table with the name 'AccountDetails'.

30. Change the name of the AccountDetails table to 'BackUpTable' table.

31. Create a view 'Acc_view' with columns CustomerId, CustomerName, AccountNumber, AccountType, and LedgerBalance from AccountsMaster. In the view Acc_view, the column names should be CustomerCode, AccountHolderName, AccountNumber, Type, and Balance for the respective columns from AccountsMaster table.

32. Create a view on AccountsMaster table with name vAccs_Dtls. This view should list all customers whose AccountType is 'IND' and their balance amount should not be less than 10000. Using this view any DML operation should not violate the view conditions.

Hint: Use the With Check Option constraint.

33. Create a view accsvw10 which will not allow DML statement against it.

34. Display the department from Staff table which has the highest salary by using Inline View.

35. List the top two highest earning employees in each department.

©2016 Capgemini. All rights reserved.

The information contained in this document is proprietary and confidential. For Capgemini only. | **19** / 29

36. Create a Sequence with the name Seq_Dept on Deptno column of Dept table. It should start from 40 and stop at 200. Increment parameter for the sequence Seq_Dept should be in step of 10.

37. Insert three sample rows by using the above sequence in Dept table.

38. Alter the above specified sequence with an increment by 5.

39. Drop the Seq_Dept sequence.

40. Create a Unique index with the name No_Name on DeptNo and Dname of Dept table.

41. Get information on the index No_Name from the Data Dictionary.

42. Create public synonym synEmp for the EMP table.

43. Get Information on synonym synEmp from the Data Dictionary.

©2016 Capgemini. All rights reserved.

The information contained in this document is proprietary and confidential. For Capgemini only. | **20** / 29

Lab 5.   Data Manipulation Language

Goals	Creating table to do the following DML operations • Insert Records • Delete Records • Update Records
Time	1 hr

5.1: Data Manipulation Language

1. Create Employee table with same structure as EMP table.

SQL>Create table employee as select * from emp where 1=3

SQL>desc employee

Name	Null?	Type
EMPNO	NOT NULL	NUMBER(4)
ENAME		VARCHAR2(10)
JOB		VARCHAR2(50)
MGR		NUMBER(4)
HIREDATE		DATE
SAL		NUMBER(7,2)
COMM		NUMBER(7,2)
DEPTNO		NUMBER(2)

SQL>select * from employee

2. Write a query to populate Employee table using EMP table's empno, ename, sal, deptno columns.

SQL>select * from employee

EMPNO	ENAME	JOB	MGR	HIREDATE	SAL	COMM	DEPTNO
7369	SMITH				800		20
7499	ALLEN				1600		30

7521	WARD				1250		30
7566	JONES				2975		20
7654	MARTIN				1250		30
7698	BLAKE				2850		30
7782	CLARK				2450		10
7788	SCOTT				3000		20
7839	KING				5000		10
7844	TURNER				1500		30
7876	ADAMS				1100		20
7900	JAMES				950		30
7902	FORD				3000		20
7934	MILLER				1300		10

14 rows selected.

3. Write a query to change the job and deptno of employee whose empno is 7698 to the job and deptno of employee having empno 7788.

4. Delete the details of department whose department name is 'SALES'.

5. Write a query to change the deptno of employee with empno 7788 to that of employee having empno 7698.

6. Insert the following rows to the Employee table through parameter substitution.

- 1000,Allen, Clerk,1001,12-jan-01, 3000, 2,10
- 1001,George, analyst, null, 08 Sep 92, 5000,0, 10
- 1002, Becker, Manager, 1000, 4 Nov 92, 2800,4, 20
- 1003, 'Bill', Clerk, 1002, 4 Nov 92,3000, 0, 20

©2016 Capgemini. All rights reserved.

The information contained in this document is proprietary and confidential. For Capgemini only. | **22** / 29

Lab 6.   Transaction Control Language Statements

Goals	• Creating a transaction • Creating a savepoint • Roll back and commit the transaction to the savepoint
Time	30 min

6.1: Transaction Control Language Statements

1. Insert rows with the following data into the Customer table. 6000, John, #115 Chicago, #115 Chicago, M, 25, 7878776, 10000

- 6001, Jack, #116 France, #116 France, M, 25, 434524, 20000

- 6002, James, #114 New York, #114 New York, M, 45, 431525, 15000.50

Use parameter substitution.

2. Create a Savepoint named 'SP1' after third record in the Customer table .

3. Insert the below row in the Customer table.

6003, John, #114 Chicago, #114 Chicago, M, 45, 439525, 19000.60

4. Execute rollback statement in such a way that whatever manipulations done before Savepoint sp1 are permanently implemented, and the ones after Savepoint SP1 are not stored as a part of the Customer table.

Appendices

Appendix A: DBMS SQL Standards

Key points to keep in mind:

NA

How to follow DBMS SQL standards:

- Decide upon a database naming convention, standardize it across your organization and be consistent in following it. It helps make your code more readable and understandable.

- Tables:

  o Tables represent the instances of an entity. For example, you store all your customer information in a table. Here "customer" is an entity and all the rows in the customers table represent the instances of the entity "customer". So name your table using the entity it represents, namely "Customer". Since the table is storing "multiple instances" of customers, make your table name a plural word.

  o Keeping this in mind:

    ▪ name your customer table as "Customers'"

    ▪ name your order storage table as "Orders"

    ▪ name your error messages table as "ErrorMessages"

  o Suppose your database deals with different logical functions and you want to group your tables according to the logical group they belong to. It will help prefixing your table name with a two or three character prefix that can identify the group.

    For example: Your database has tables which store information about Sales and Human resource departments, then you can name all your tables related to Sales department as shown below:

      ▪ SL_NewLeads

      ▪ SL_Territories

      ▪ SL_TerritoriesManagers

    You can name all your tables related to Human resources department as shown below:

      ▪ HR_Candidates

      ▪ HR_PremierInstitutes

      ▪ HR_InterviewSchedules

- Prefix the table names with owner names, as this improves readability, and avoids any unnecessary confusion.

- Primary keys:
  - Primary key is the column(s) that can uniquely identify each row in a table. So just use the column name prefixed with "pk_" + "Table name" for naming primary keys.
  - Given below is an example of how we can name the primary key on the CustomerID column of Customers table:
    - pk_Customers_CustomerID
  - Consider concatenating the column names in case of composite primary keys.
- Foreign keys:
  - Foreign keys are used to represent the relationships between tables which are related. So a foreign key can be considered as a link between the "column of a referencing table" and the "primary key column of the referenced table".
  - We can use the following naming convention for foreign keys:
    - fk_referencing table + referencing column_referenced table + referenced column
  - Based on the above convention, we can name the foreign key, which references the CustomerID column of the Customers table from the Order's tables CustomerID column as shown below:
    - fk_OrdersCustomerID_CustomersCustomerID
  - Foreign key can be composite too. In that case, consider concatenating the column names of referencing and referenced tables while naming the foreign key. This might make the name of the foreign key lengthy. However, you should not be worried about it, as you will never reference this name from your code, except while creating/dropping these constraints.
- Default and Check constraints:
  - Use the column name to which the defaults /check constraints are bound to. Prefix it with "def" and "chk" prefixes respectively for Default and Check constraints.
  - We can name the default constraint for OrderDate Column as def_OrderDate, and the check constraint for OrderDate column as chk_OrderDate.
- Do not use any reserved words for naming my database objects, as that can lead to some unpredictable situations.
- Do not depend on undocumented functionality. The reasons are:

- o you will not get support, when something goes wrong with your undocumented code

- o undocumented functionality is not guaranteed to exist (or behave the same) in a future release or service pack, thereby breaking your code

- Make sure you normalize your data at least till third normal form. At the same time, do not compromise on query performance. A little de-normalization helps queries perform faster.

Appendix B: Coding Best Practices

Given below are a few best practices in DBMS SQL:

- Do not use SELECT * in your queries. Always write the required column names after the SELECT statement, as shown below:
      SELECT CustomerID, CustomerFirstName, City

    This technique results in less disk IO, less network traffic, and hence better performance.

- Try to avoid wildcard characters at the beginning of a word while searching by using the LIKE keyword. This is because this arrangement results in an index scan, which is defeating the purpose of having an index. The following statement results in an index scan, while the second statement results in an index seek:

      1. SELECT LocationID FROM Locations WHERE Specialities LIKE '%pples'
      2. SELECT LocationID FROM Locations WHERE Specialities LIKE 'A%s'

    Also avoid searching with not equals operators (<> and NOT) as they result in table and index scans. If you must do heavy text-based searches, consider using the Full-Text search feature of SQL Server for better performance.

- Use "Derived tables" wherever possible, as they perform better. Consider the following query to find the second highest salary from Employees table:

```
SELECT MIN(Salary)
FROM Employees
WHERE EmpID IN
(
SELECT TOP 2 EmpID
FROM Employees
ORDER BY Salary Desc
)
```

    The same query can be re-written by using a derived table as shown below, and it performs twice as fast as the above query:

```
SELECT MIN(Salary)
FROM
(
SELECT TOP 2 Salary
FROM Employees
ORDER BY Salary Desc
) AS A
```

This is just an example. The results might differ in different scenarios depending upon the database design, indexes, volume of data, etc. So test all the possible ways a query can be written and go with the efficient one.

- Use char data type for a column, only when the column is non-nullable. If a char column is nullable, it is treated as a fixed length column in SQL Server 7.0+. So a char(100), when NULL, will eat up 100 bytes, resulting in space wastage. So use varchar(100) in this situation. Of course, variable length columns do have a very little processing overhead over fixed length columns. Carefully choose between char and varchar depending upon the length of the data you are going to store.

- Always use a column list in your INSERT statements. This helps in avoiding problems when the table structure changes (like adding a column).

- Do not use the column numbers in the ORDER BY clause as it impairs the readability of the SQL statement. Further, changing the order of columns in the SELECT list has no impact on the ORDER BY when the columns are referred by names instead of numbers. Consider the following example, in which the second query is more readable than the first one:

SELECT OrderID, OrderDate
FROM Orders
ORDER BY 2

SELECT OrderID, OrderDate
FROM Orders
ORDER BY OrderDate

Appendix C: Table of Examples